



# Deep Learning for ECE

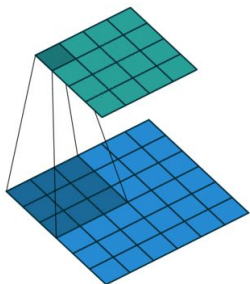
## EECE-580G

Advanced CNNs

# Last time

# CNNs ingredients

---

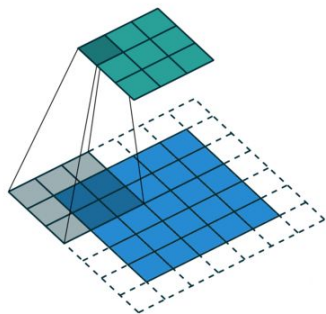


conv2d

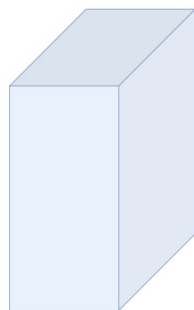
0	0	0	0	0	0	0	0
0	3	3	2	1	0	0	0
0	0	0	1	3	1	0	0
0	3	1	2	2	3	0	0
0	2	0	0	2	2	0	0
0	2	0	0	0	1	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0

pool with stride

3.0	3.0	3.0
3.0	3.0	3.0
2.0	2.0	2.0



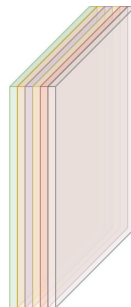
conv2d with stride



$A^{[i-1]}$



$K^{[i]}$



$Z^{[i]}$

conv2d on volumes

## Batch Norm update (one layer)

---

$B$  minibatch of size  $m$ ,  $B = \{x^{(1)}, \dots, x^{(m)}\}$ ,  $x^{(i)} \in \mathbb{R}^d$

$$\mu^{(B)} = \frac{1}{m} \sum_{i=1}^m x^{(i)} \quad \sigma^{(B)2} = \frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu^{(B)})^2$$

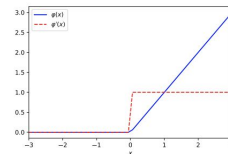
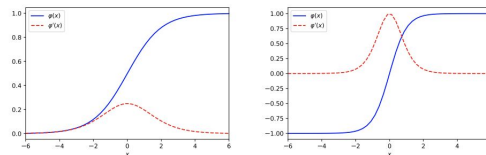
normalize

$$\hat{x}_k^{(i)} = \frac{x_k^{(i)} - \mu_k^{(B)}}{\sqrt{\sigma_k^{(B)2} + \epsilon}}$$

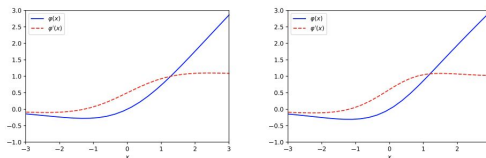
output

$$y_k^{(i)} = \gamma_k \hat{x}_k^{(i)} + \beta_k \quad \gamma, \beta \text{ trainable variables}$$

## Batch Norm



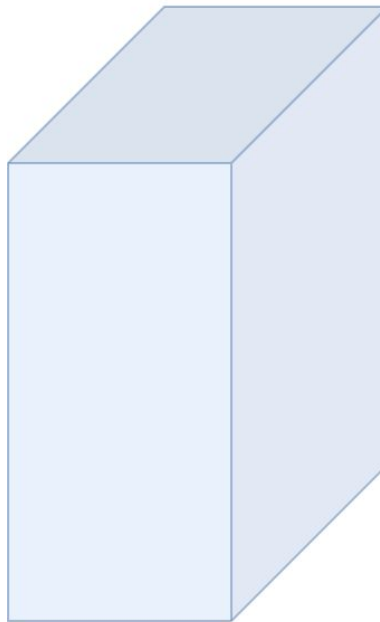
## Activation functions



# 1x1 convolution

— — —

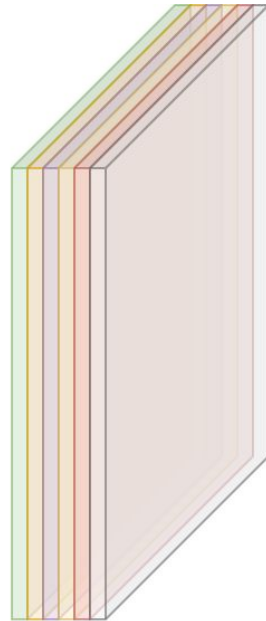
- Does not mix spatial information
- Only mixes channels
- Each 1x1 conv return a weighted avg of the input channels
- Useful when need to decrease/increase the depth independently of the spatial info
- Can you see that a 1x1 conv on a 1x1xD image is equivalent to flatten  $\Rightarrow$  FC?  $\gg \gg$  **Piazza**



$A^{[i-1]}$



$K^{[i]}$



$Z^{[i]}$

# The “conv” layer

— — —

- Usually = **conv - BN - activation**
- When followed by BN no need to use the bias in conv layer

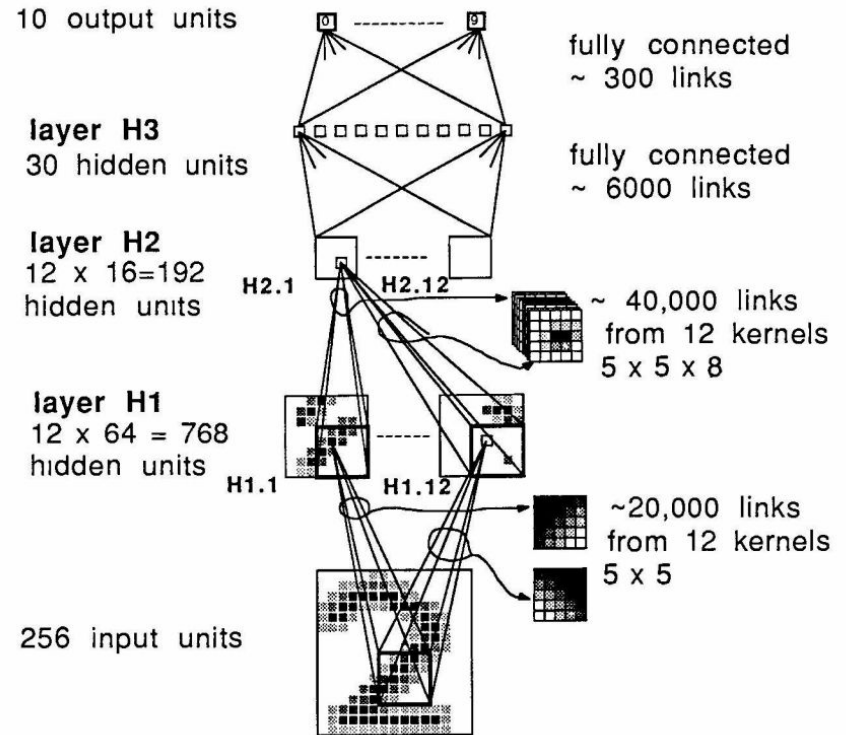
# Early CNNs

# First CNNs (1989)

- Yann LeCun - Bell Labs
- Trained for 3 weeks
- Activation = tanh
- Strided convolutions at each layer

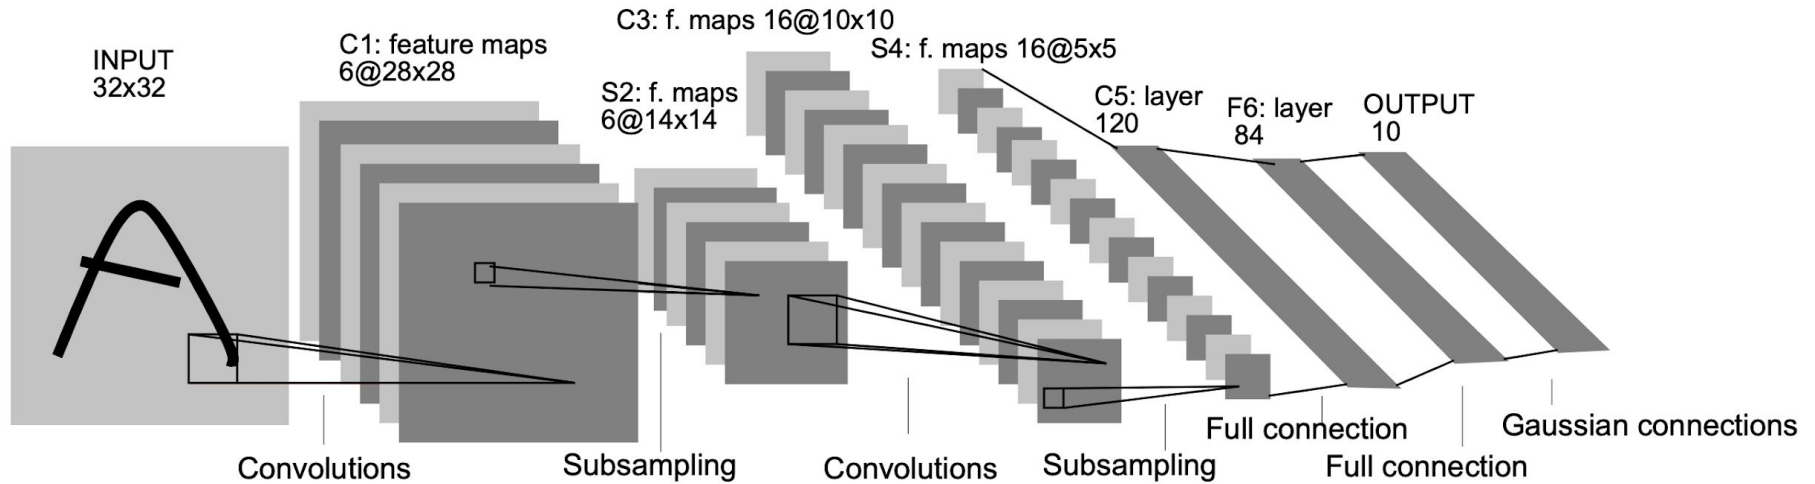


Source: [NYU.edu](https://www.nyu.edu)



[LeCun, Yann, et al. "Backpropagation applied to handwritten zip code recognition." Neural computation 1.4 \(1989\): 541-551.](#)

# LeNet-5 (1998)



[LeCun, Yann, et al. "Gradient-based learning applied to document recognition." Proceedings of the IEEE 86.11 \(1998\): 2278-2324.](#)



# Missing Ingredient: Hardware...

— — —

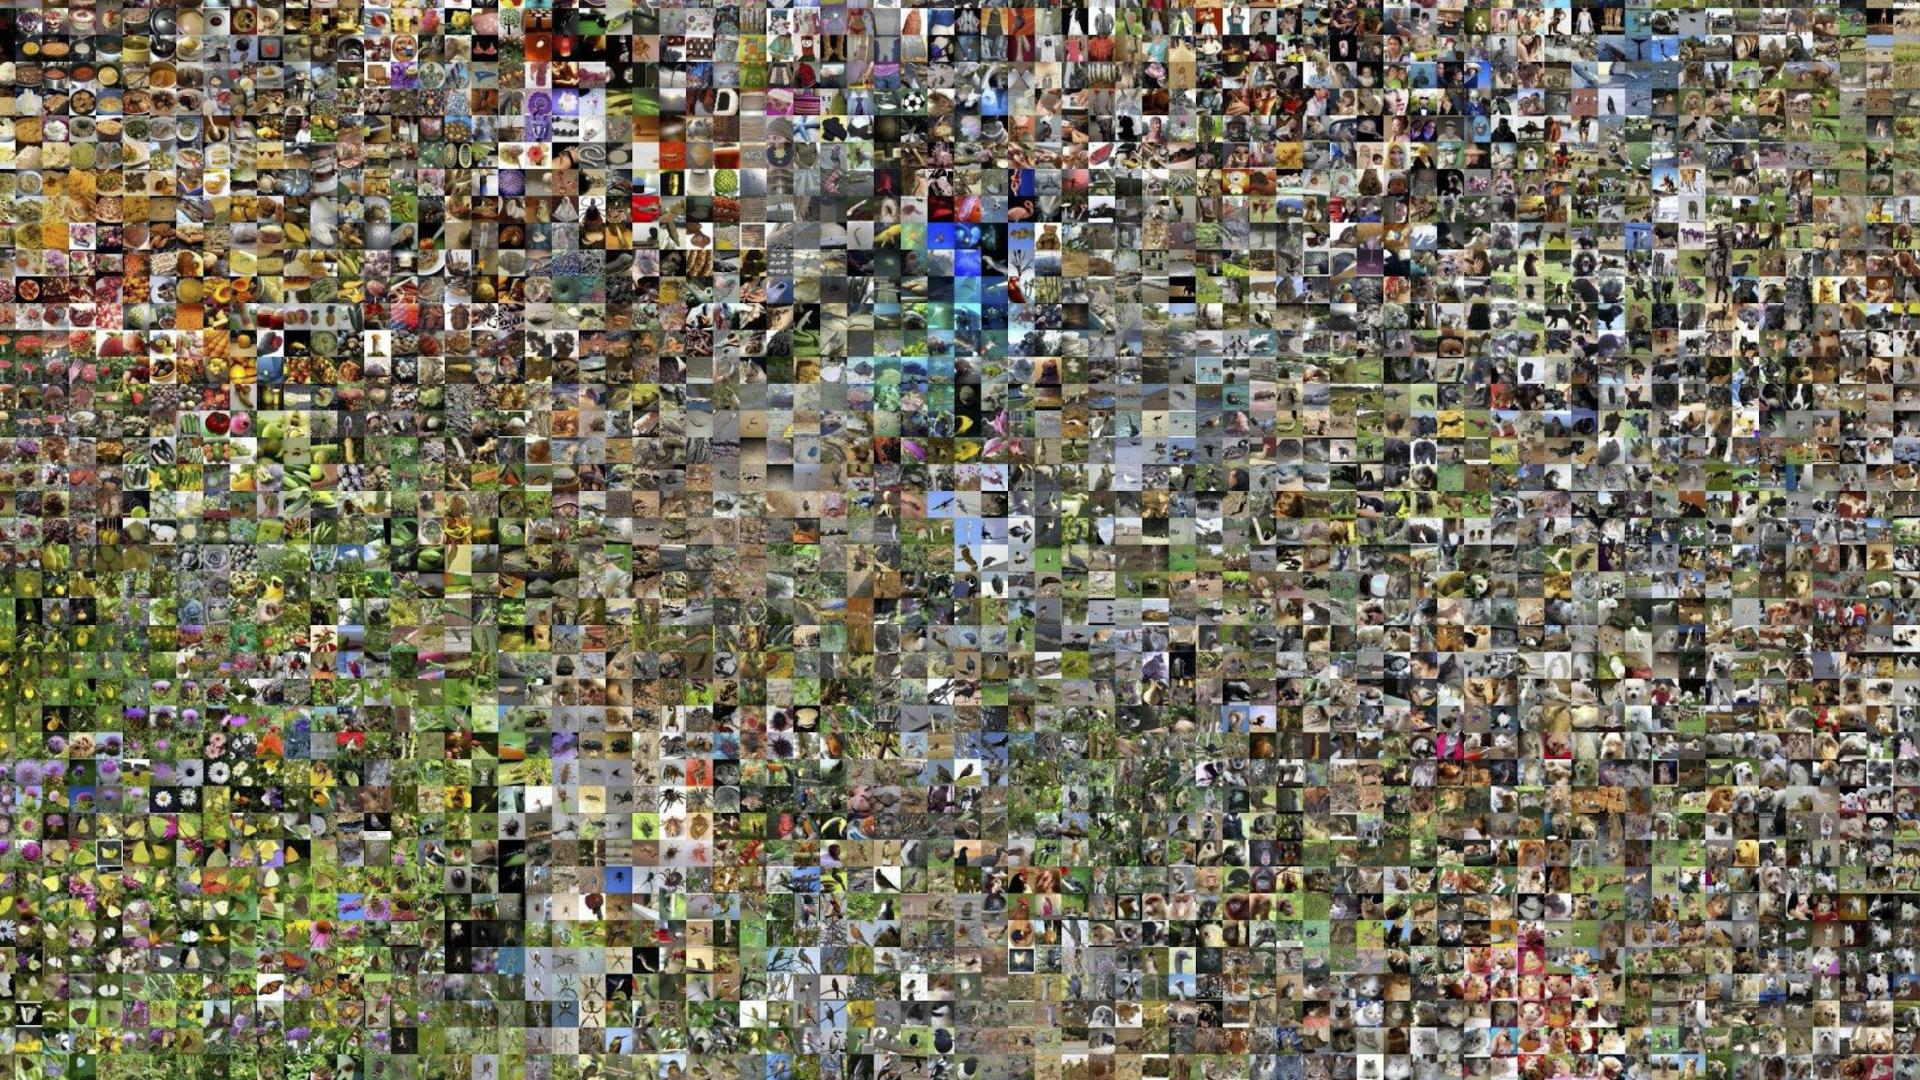
- DL research slowed down due to inefficient CPU training ...
- **Early 2000's** = GPU development thanks to gaming industry



2012 GPU GeForce GTX 580

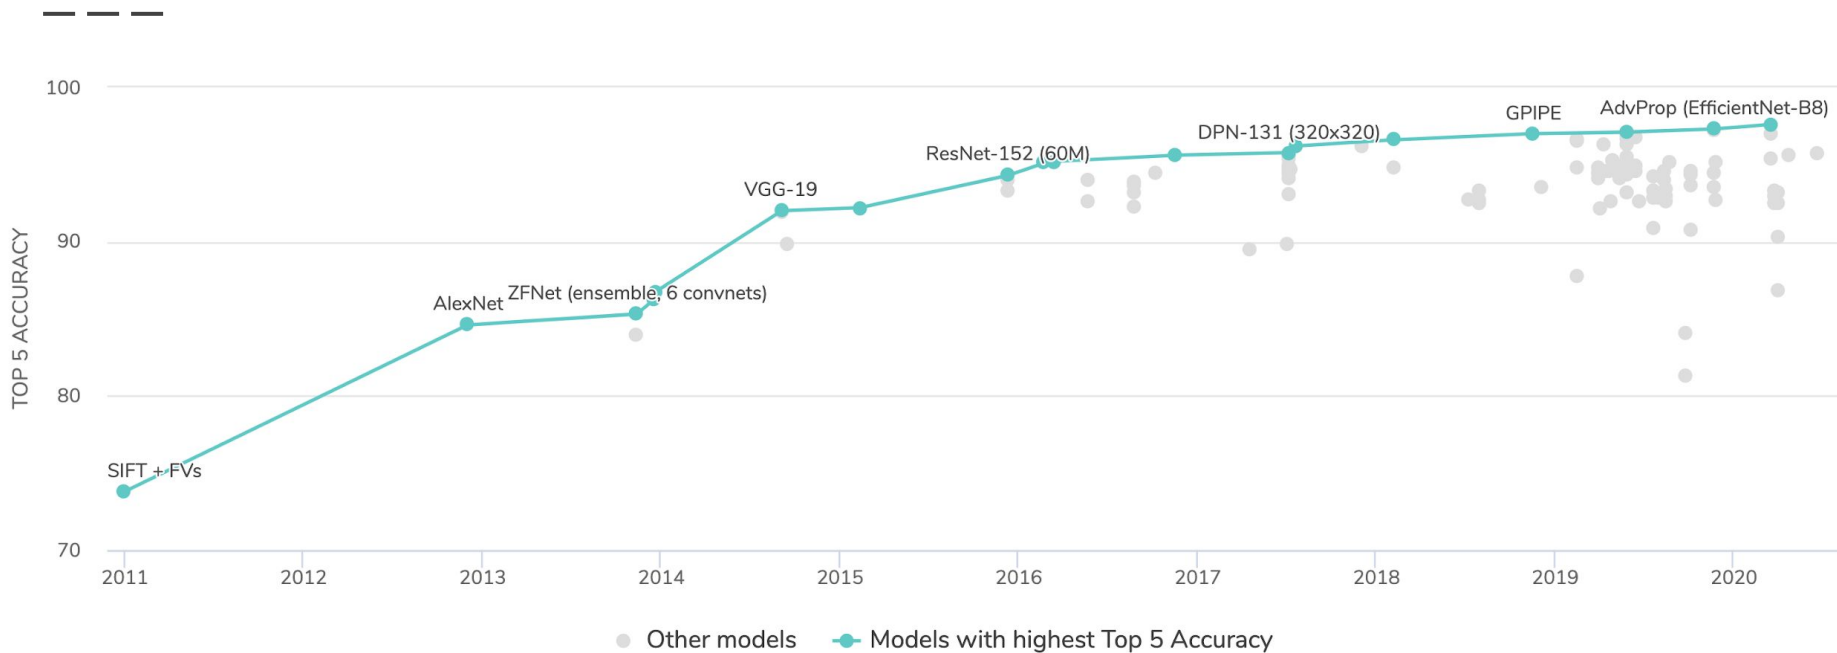
# The ImageNet Challenge





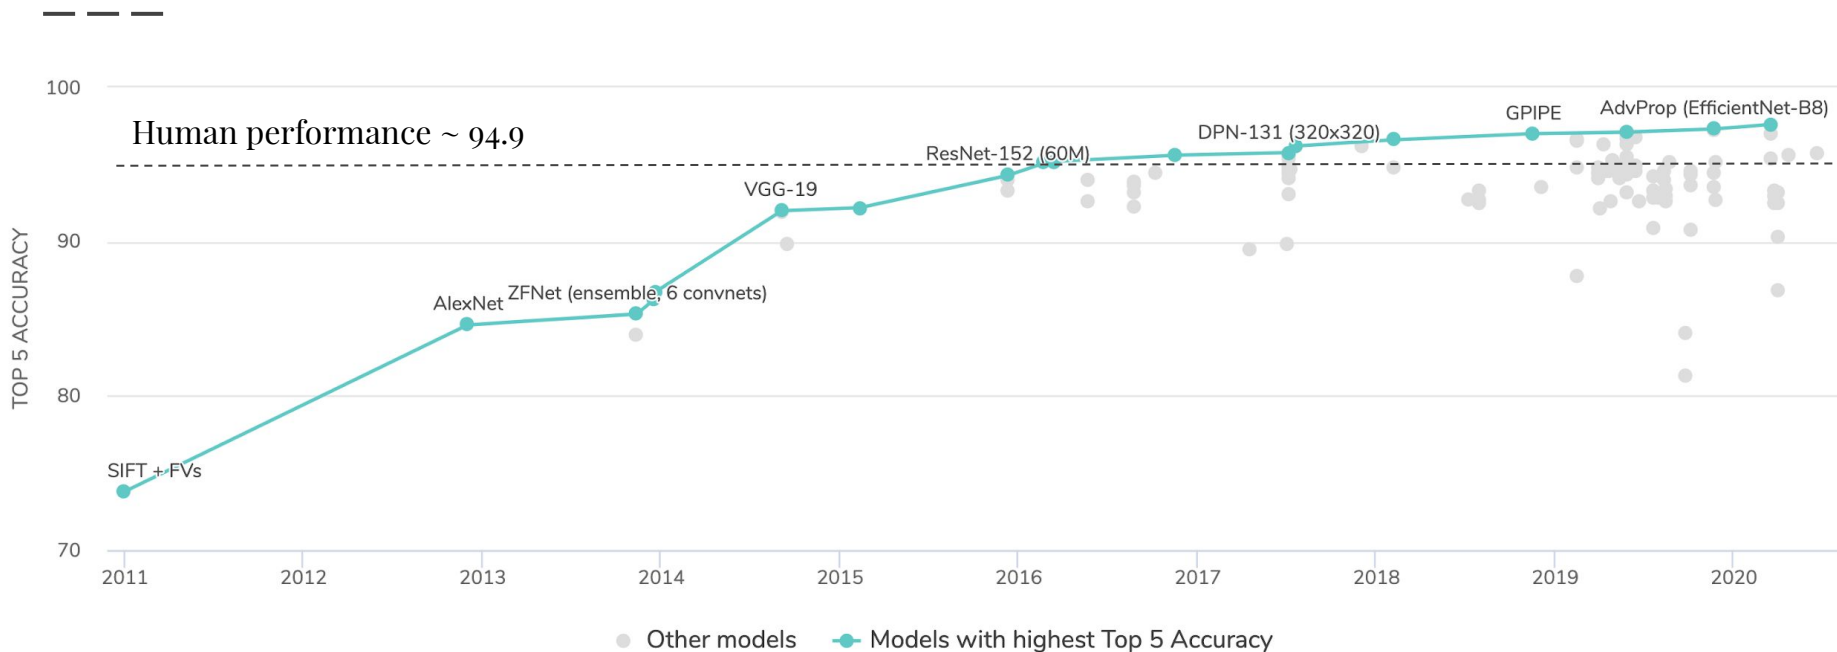


# ImageNet Classification Challenge



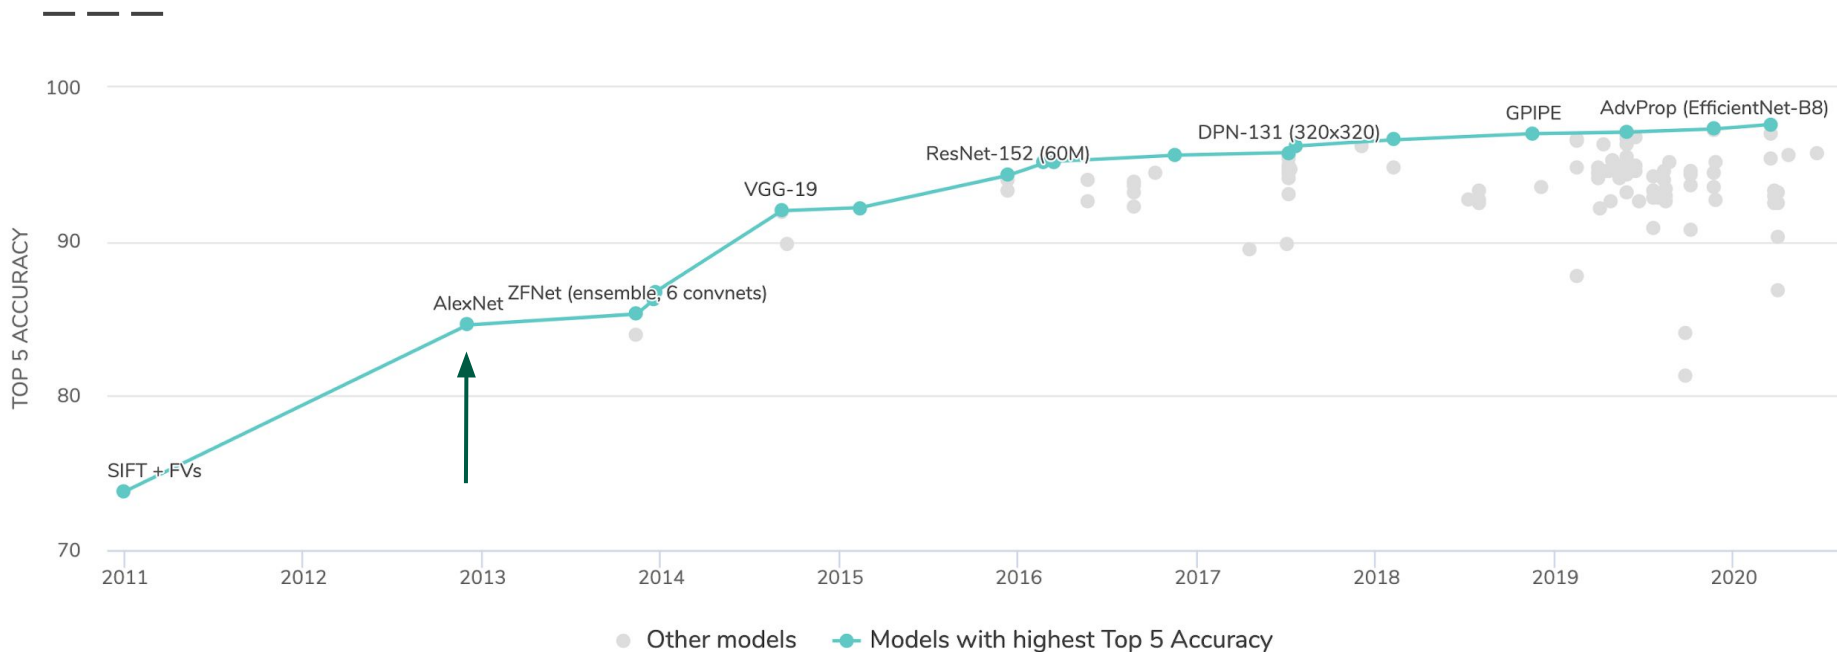
<https://paperswithcode.com/sota/image-classification-on-imagenet>

# ImageNet Classification Challenge



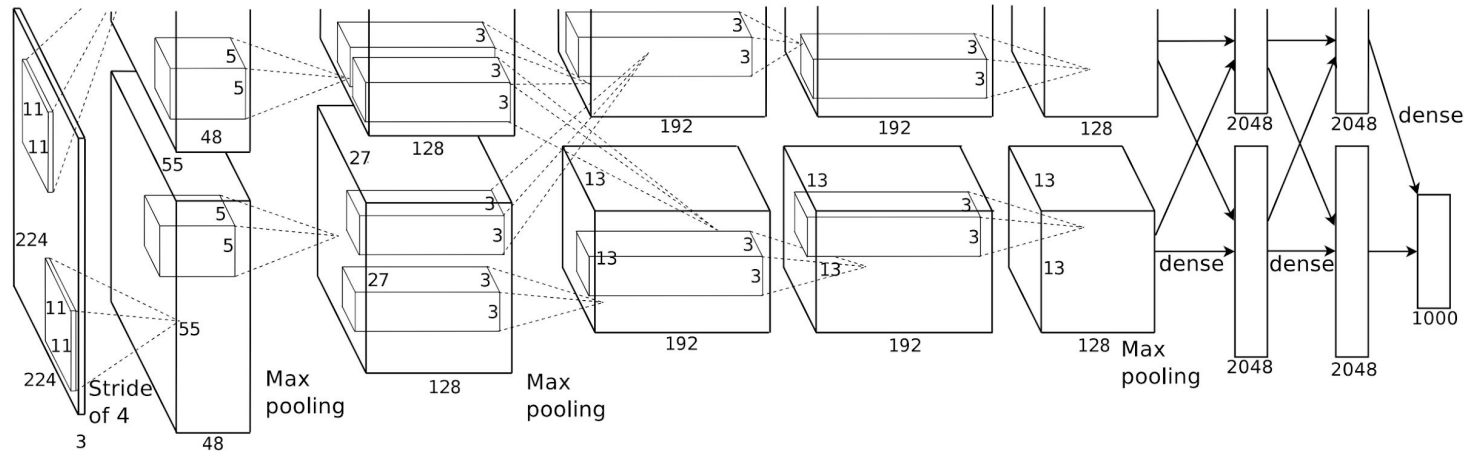
<https://paperswithcode.com/sota/image-classification-on-imagenet>

# ImageNet Classification Challenge



<https://paperswithcode.com/sota/image-classification-on-imagenet>

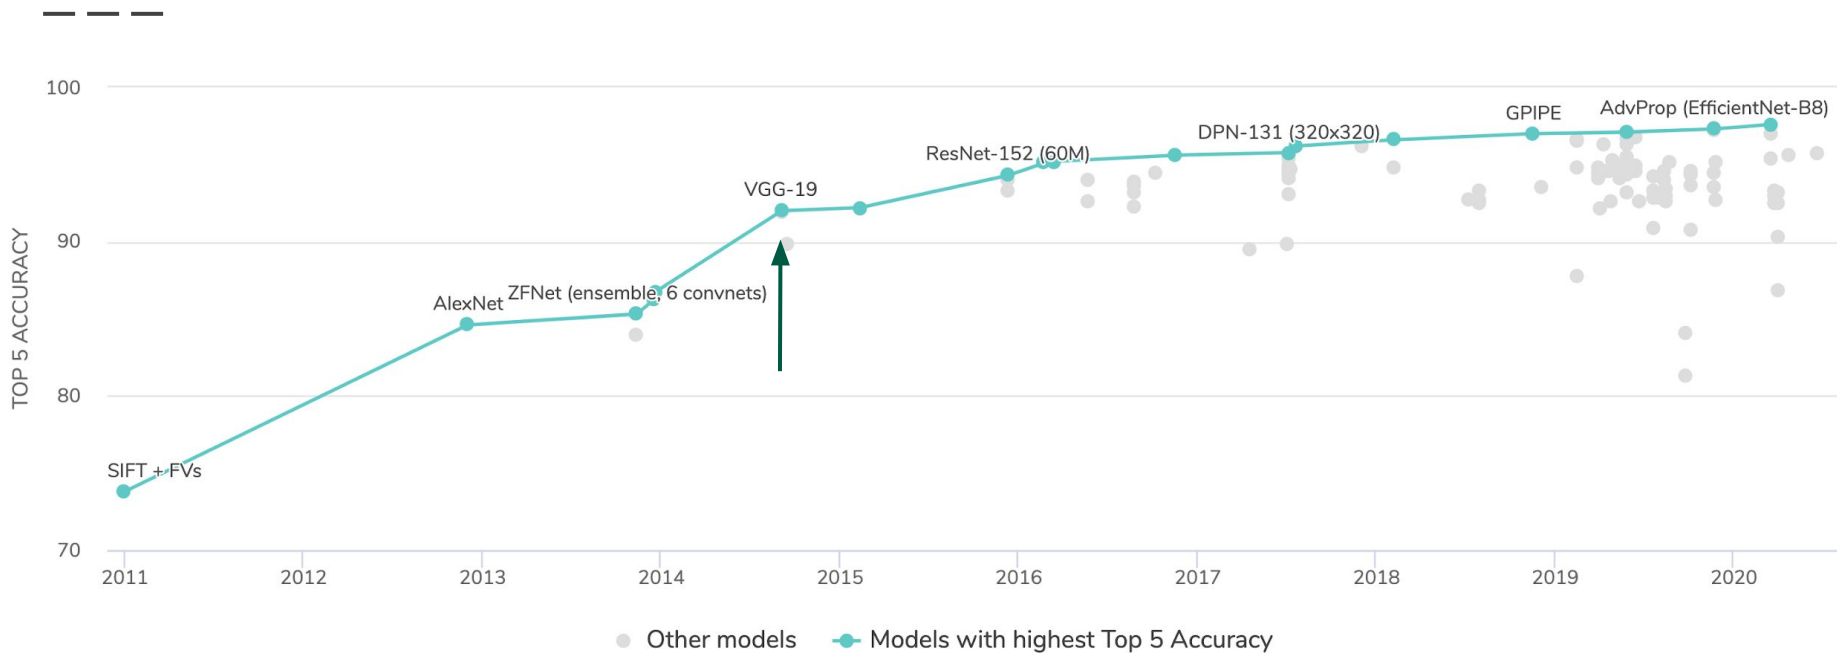
# AlexNet



[Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E. Hinton. "Imagenet classification with deep convolutional neural networks." Advances in neural information processing systems. 2012](#)

- Local Response Normalization (Similar to Batch Norm in principle, Batch Norm was introduced in 2015)
- Cited 71,641 times!
- ReLU activation

# ImageNet Classification Challenge



<https://paperswithcode.com/sota/image-classification-on-imagenet>



# VGG

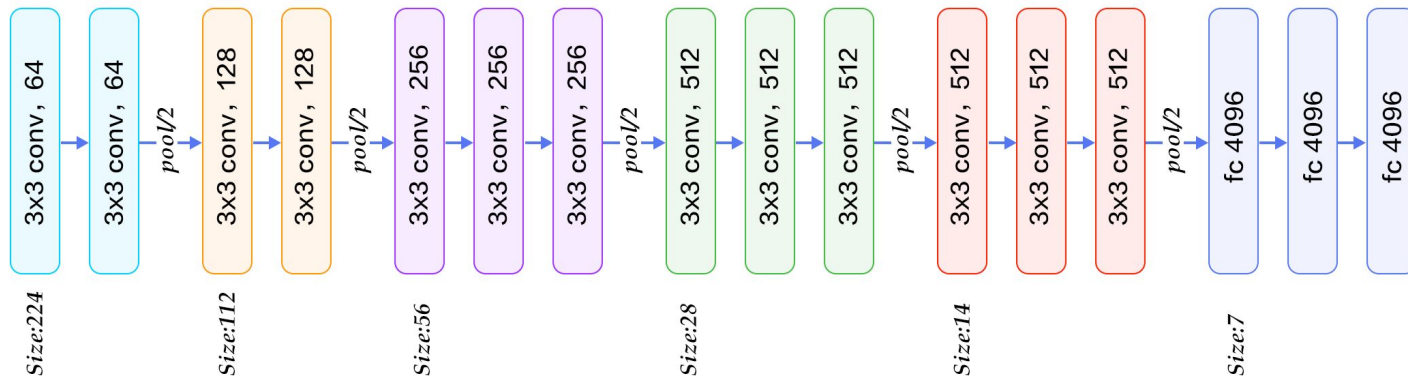


Figure Credits: Yugandhar Nanda's answer on Quora: [What is the VGG neural network?](#)

VGG paper: [Simonyan, K. and A. Zisserman. "Very Deep Convolutional Networks for Large-Scale Image Recognition." CoRR abs/1409.1556 \(2015\)](#)

- Local Response Normalization not used
- Fix kernel size to 3x3 – increase receptive field by stacking more 3x3 conv
- ReLU activation

(Not including biases for simplicity)

**INPUT:** [224x224x3] memory:  $224*224*3=150\text{K}$  params: 0

**CONV3-64:** [224x224x64] memory:  $224*224*64=3.2\text{M}$  params:  $(3*3*3)*64 = 1,728$

**CONV3-64:** [224x224x64] memory:  $224*224*64=3.2\text{M}$  params:  $(3*3*64)*64 = 36,864$

**POOL2:** [112x112x64] memory:  $112*112*64=800\text{K}$  params: 0

**CONV3-128:** [112x112x128] memory:  $112*112*128=1.6\text{M}$  params:  $(3*3*64)*128 = 73,728$

**CONV3-128:** [112x112x128] memory:  $112*112*128=1.6\text{M}$  params:  $(3*3*128)*128 = 147,456$

**POOL2:** [56x56x128] memory:  $56*56*128=400\text{K}$  params: 0

**CONV3-256:** [56x56x256] memory:  $56*56*256=800\text{K}$  params:  $(3*3*128)*256 = 294,912$

**CONV3-256:** [56x56x256] memory:  $56*56*256=800\text{K}$  params:  $(3*3*256)*256 = 589,824$

**CONV3-256:** [56x56x256] memory:  $56*56*256=800\text{K}$  params:  $(3*3*256)*256 = 589,824$

**POOL2:** [28x28x256] memory:  $28*28*256=200\text{K}$  params: 0

**CONV3-512:** [28x28x512] memory:  $28*28*512=400\text{K}$  params:  $(3*3*256)*512 = 1,179,648$

**CONV3-512:** [28x28x512] memory:  $28*28*512=400\text{K}$  params:  $(3*3*512)*512 = 2,359,296$

**CONV3-512:** [28x28x512] memory:  $28*28*512=400\text{K}$  params:  $(3*3*512)*512 = 2,359,296$

**POOL2:** [14x14x512] memory:  $14*14*512=100\text{K}$  params: 0

**CONV3-512:** [14x14x512] memory:  $14*14*512=100\text{K}$  params:  $(3*3*512)*512 = 2,359,296$

**CONV3-512:** [14x14x512] memory:  $14*14*512=100\text{K}$  params:  $(3*3*512)*512 = 2,359,296$

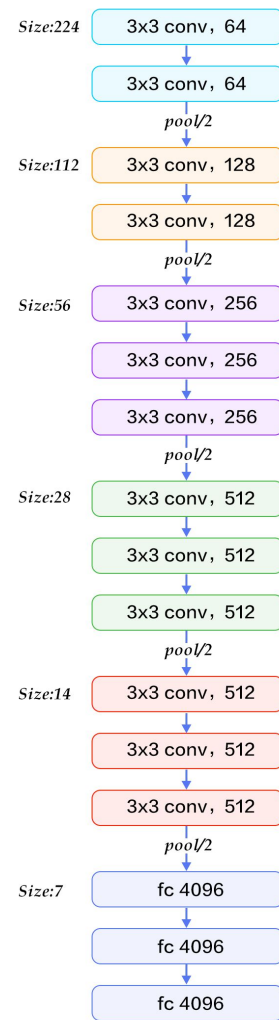
**CONV3-512:** [14x14x512] memory:  $14*14*512=100\text{K}$  params:  $(3*3*512)*512 = 2,359,296$

**POOL2:** [7x7x512] memory:  $7*7*512=25\text{K}$  params: 0

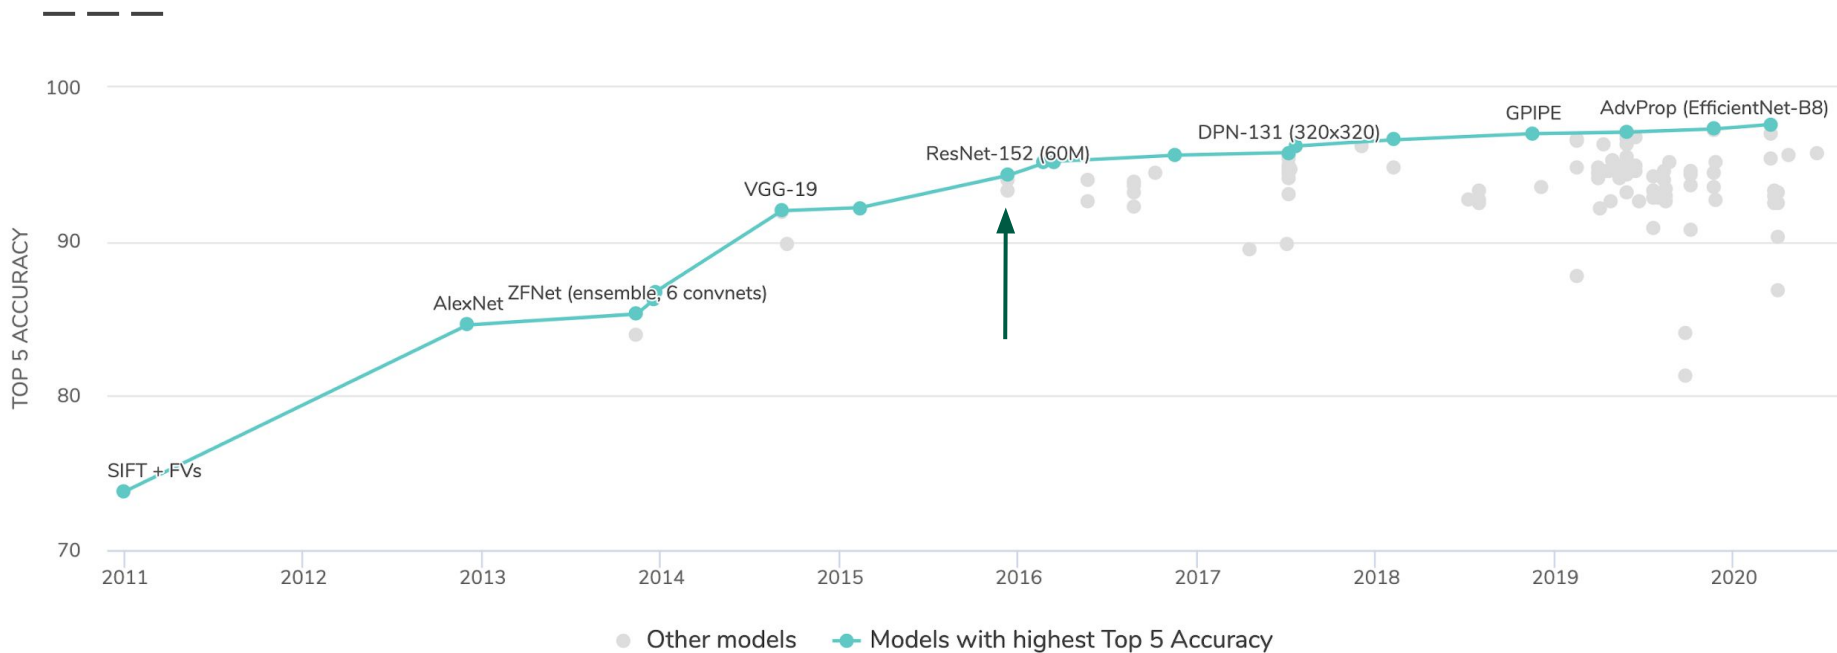
**FC:** [1x1x4096] memory: 4096 params:  $7*7*512*4096 = 102,760,448$

**FC:** [1x1x4096] memory: 4096 params:  $4096*4096 = 16,777,216$

**FC:** [1x1x1000] memory: 1000 params:  $4096*1000 = 4,096,000$



# ImageNet Classification Challenge



<https://paperswithcode.com/sota/image-classification-on-imagenet>

# ResNet

# He, Kaiming, et al. "Deep residual learning for image recognition." CVPR 2016.

— — —

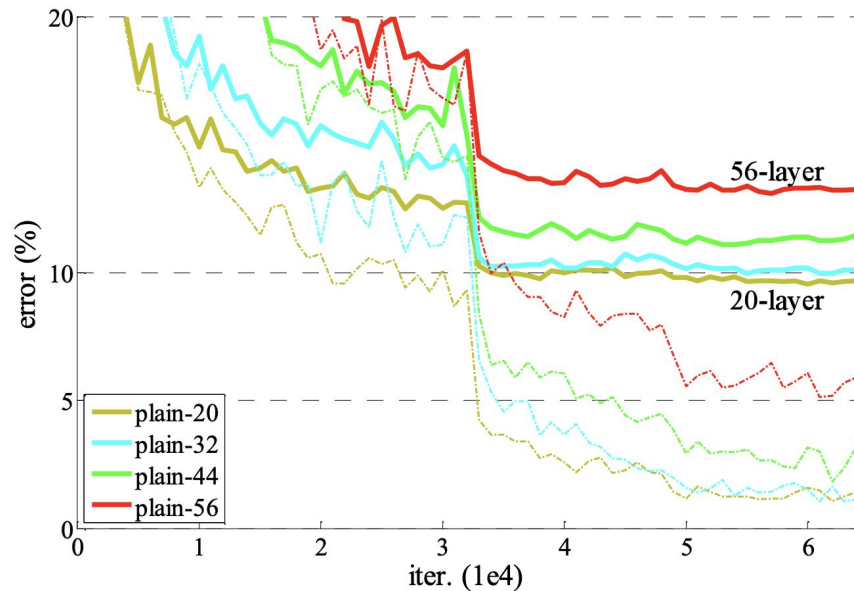
## Abstract

*Deeper* neural networks are **more difficult to train**. We present a residual learning framework to ease the training of networks that are substantially deeper than those used previously. We explicitly reformulate the layers as learning residual functions with reference to the layer inputs, instead of learning unreferenced functions. We provide comprehensive empirical evidence showing that these residual networks are easier to optimize, and can **gain accuracy from considerably increased depth**. On the ImageNet dataset we evaluate residual nets with a depth of up to 152 layers— $8\times$  deeper than VGG nets but still having lower complexity. An ensemble of these residual nets achieves 3.57% error on the ImageNet test set. This result won the 1st place on the ILSVRC 2015 classification task. We also present analysis on CIFAR-10 with 100 and 1000 layers.

Credits: The following slides use figures from the ResNet paper

# ResNet

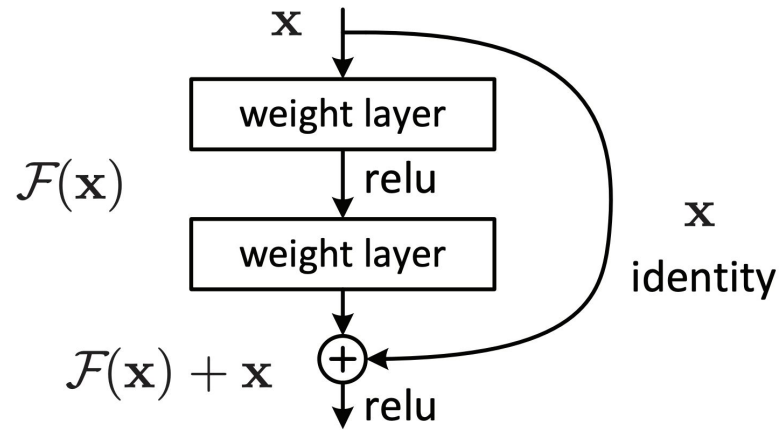
- **Deeper** CNNs give **worse** performance?
- Hard to train deeper CNNs
- Not an overfitting problem
- In theory  
A deeper CNN should be able to represent a shallower CNN?
- Not easy in practice...



# Residual blocks

— — —

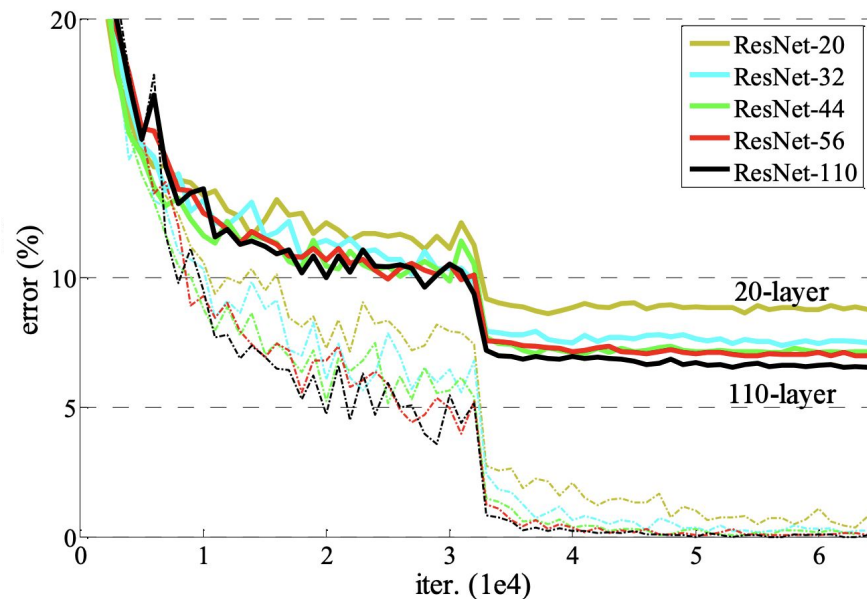
- Learn the residual
- Suppose you want to learn:  $\mathcal{H}(x)$
- Residual blocks learn:  $\mathcal{F}(x) = \mathcal{H}(x) - x$
- Output is equivalent
- The identity connection is called **shortcut** or **skip** connection
- **Easier to learn a zero mapping than an identity mapping**



# ResNet

— — —

- Easier to train residual CNNs
- **Deeper = Better**
- Residual blocks = the new norms in conv blocks

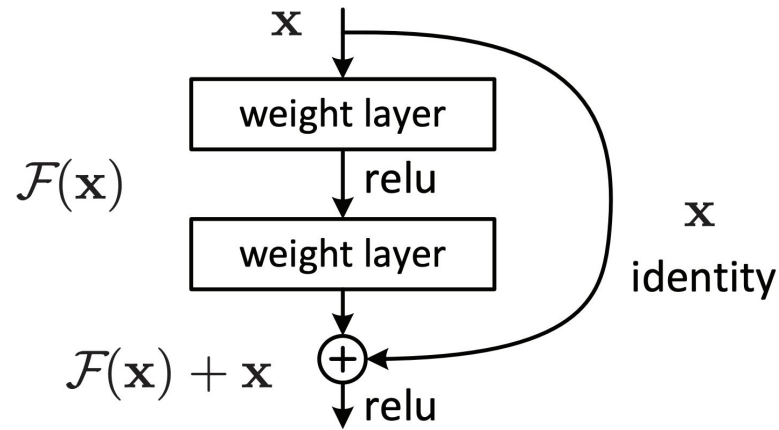




# Residual blocks

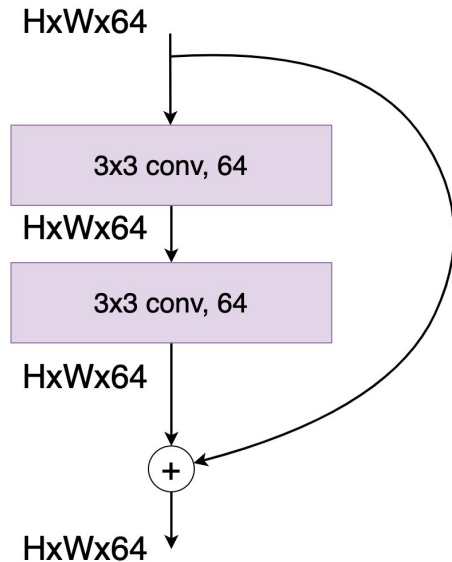
— — —

- Learn the residual
- Suppose you want to learn:  $\mathcal{H}(x)$
- Residual blocks learn:  $\mathcal{F}(x) = \mathcal{H}(x) - x$
- Output is equivalent
- The identity connection is called **shortcut** or **skip** connection
- **Easier to learn a zero mapping than an identity mapping**
- **Can we always add the input back to the output of conv operations?**



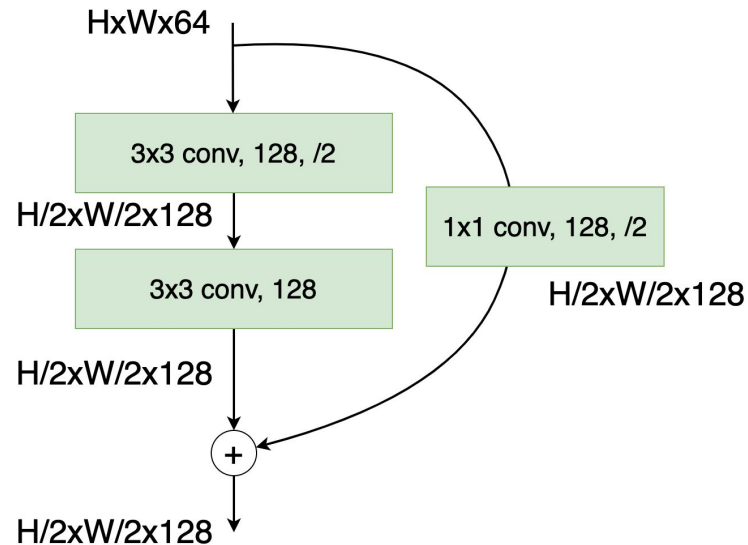
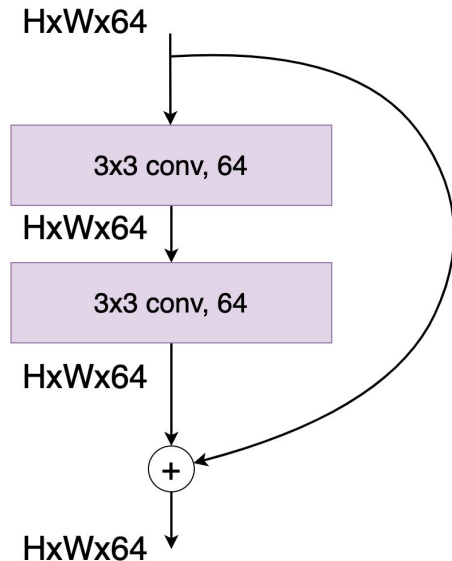
# Residual blocks

- If the output of the main branch has smaller resolution (due to stride or pooling)
- The shortcut connection has to bring  $x$  to that resolution
- Using  $1\times 1$  conv-stride



# Residual blocks

- If the output of the main branch has smaller resolution (due to stride or pooling)
- The shortcut connection has to bring  $x$  to that resolution
- Using  $1\times 1$  conv-stride

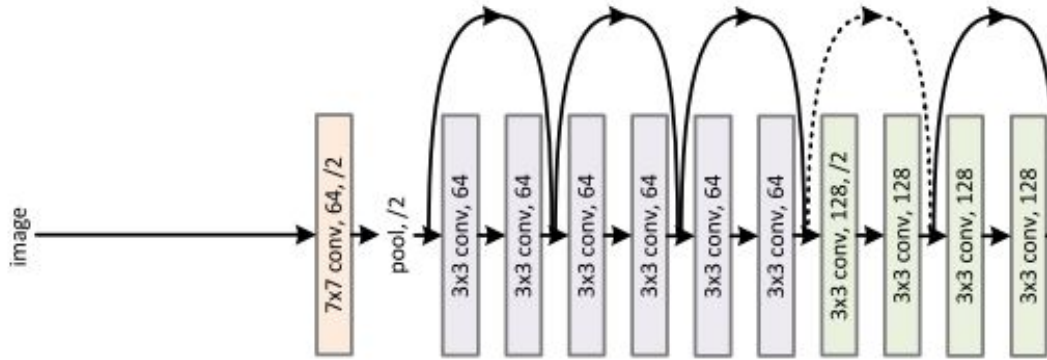


\_\_\_\_\_

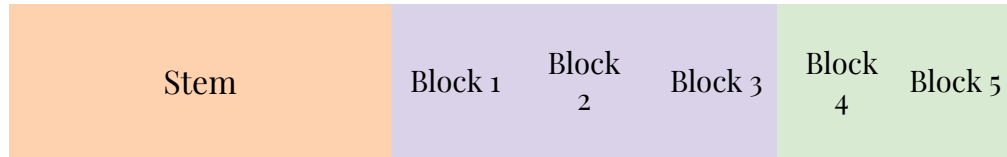
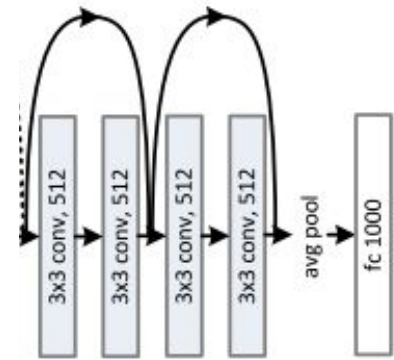


# Terminology

34-layer residual

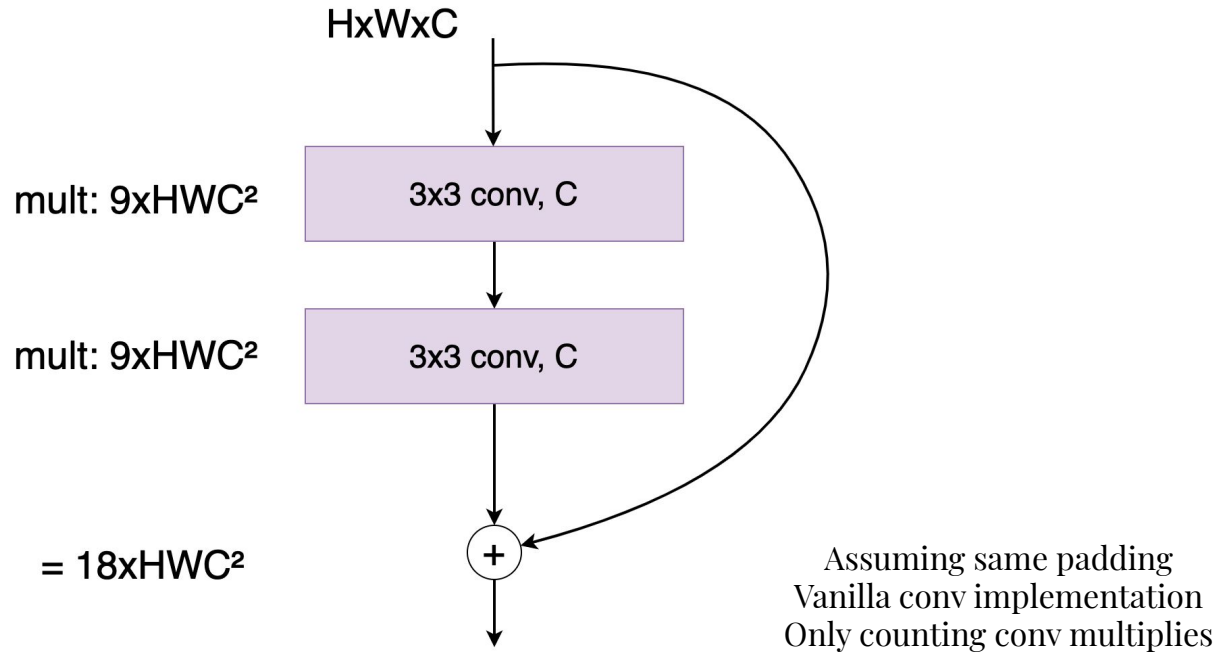


...



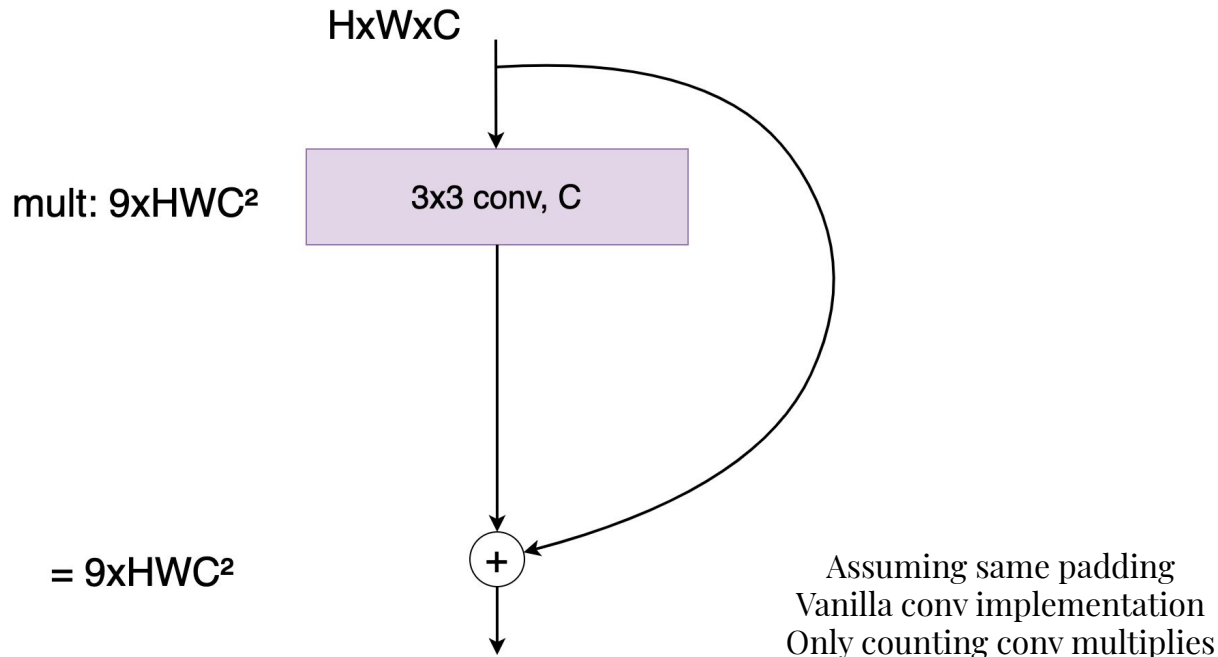
# Going deeper with the bottleneck block

— — —

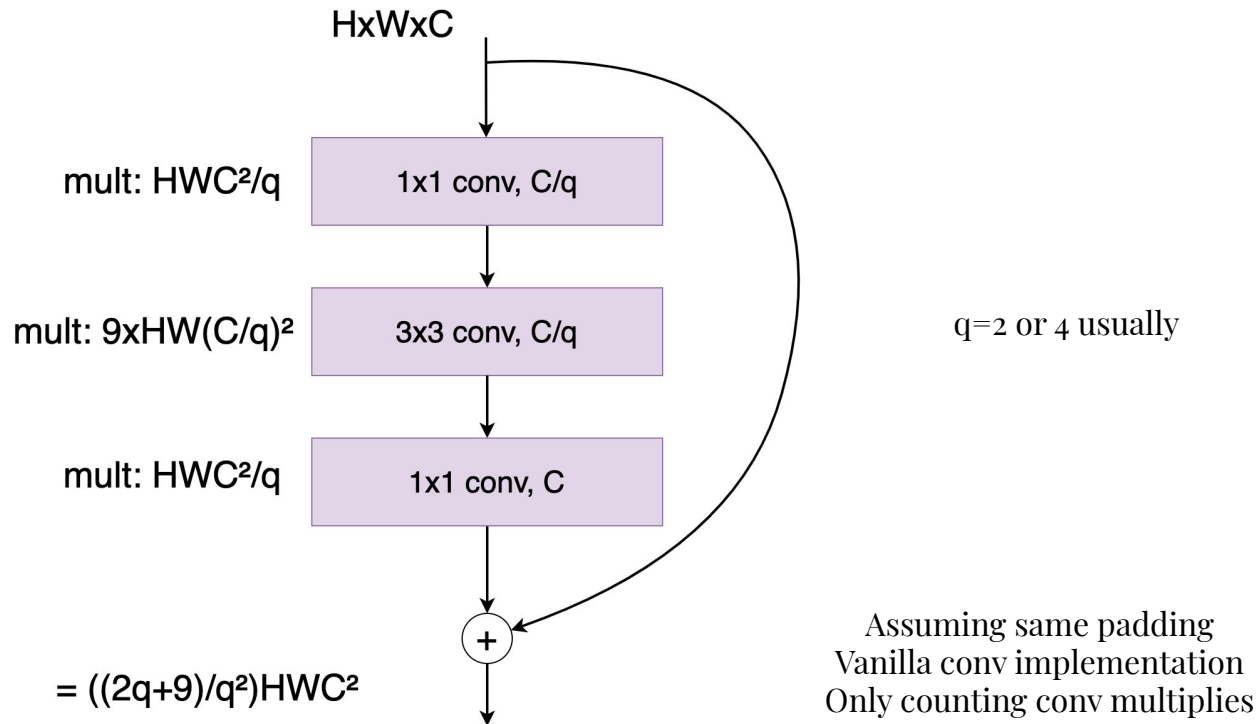


# Going deeper with the bottleneck block

— — —



# Going deeper with the bottleneck block





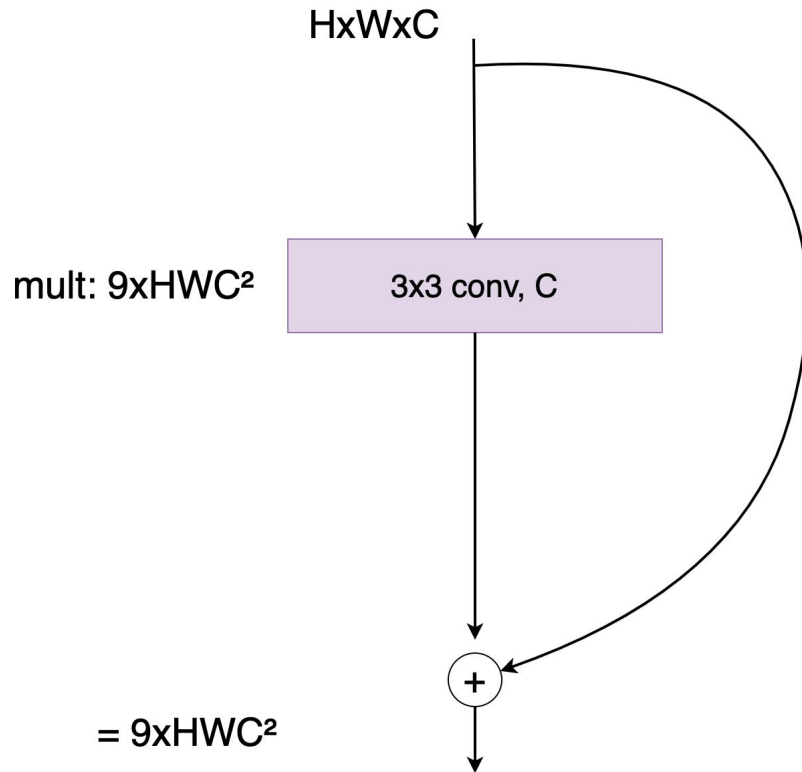
# Going deeper with the bottleneck block

— — —

Architecture	GFLOPS	ImageNet top 5 error
ResNet-34 (basic res block)	3.6	8.58
ResNet-50 (bottleneck block)	3.8	7.13
ResNet-152 (bottleneck block)	11.3	5.84

# Going deeper with the MobileNet block

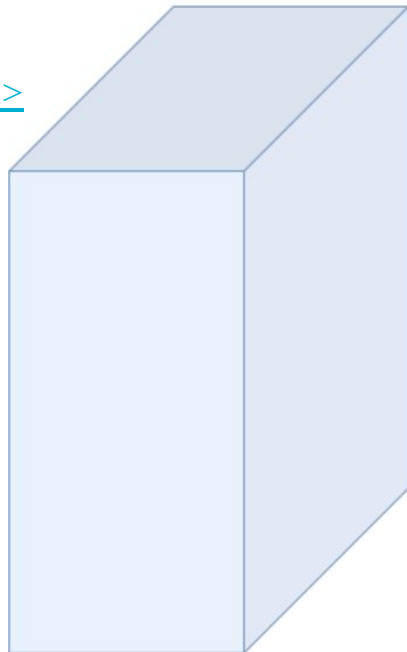
— — —



Assuming same padding  
Vanilla conv implementation  
Only counting conv multiplies

# Depthwise (separable) convolution

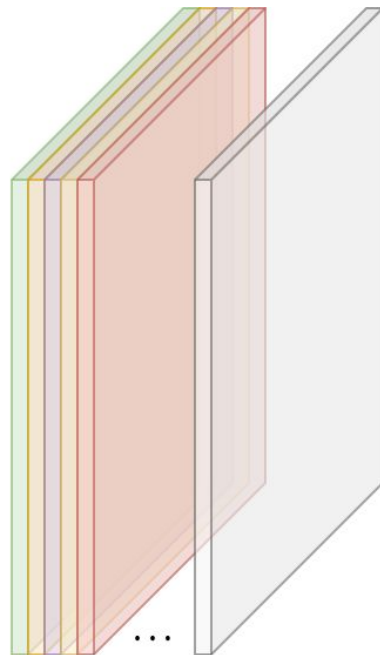
Implemented as >>>



$A^{[i-1]}$

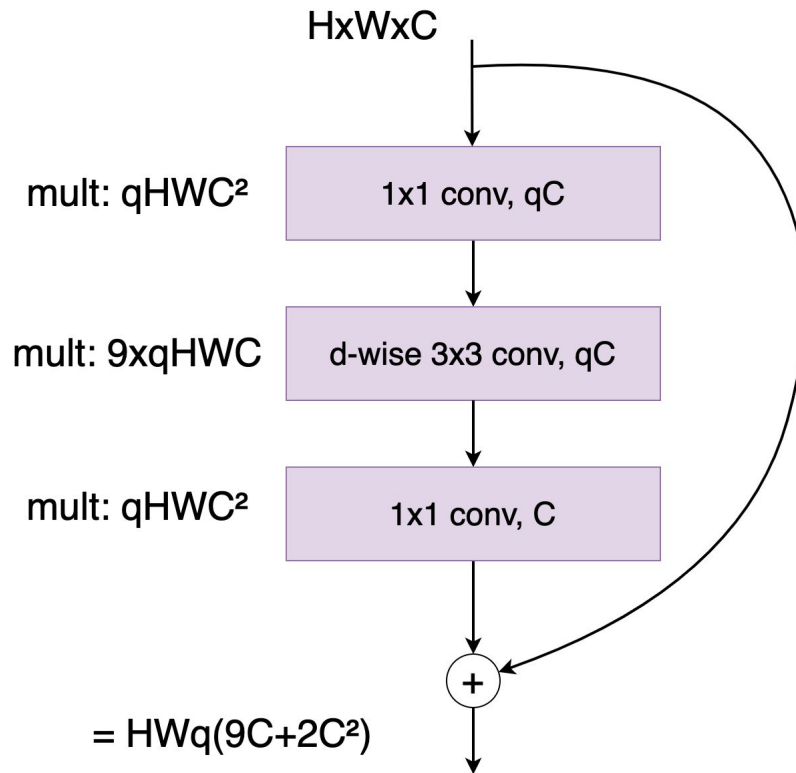


$K^{[i]}$



$Z^{[i]}$

# Going deeper with the MobileNet block



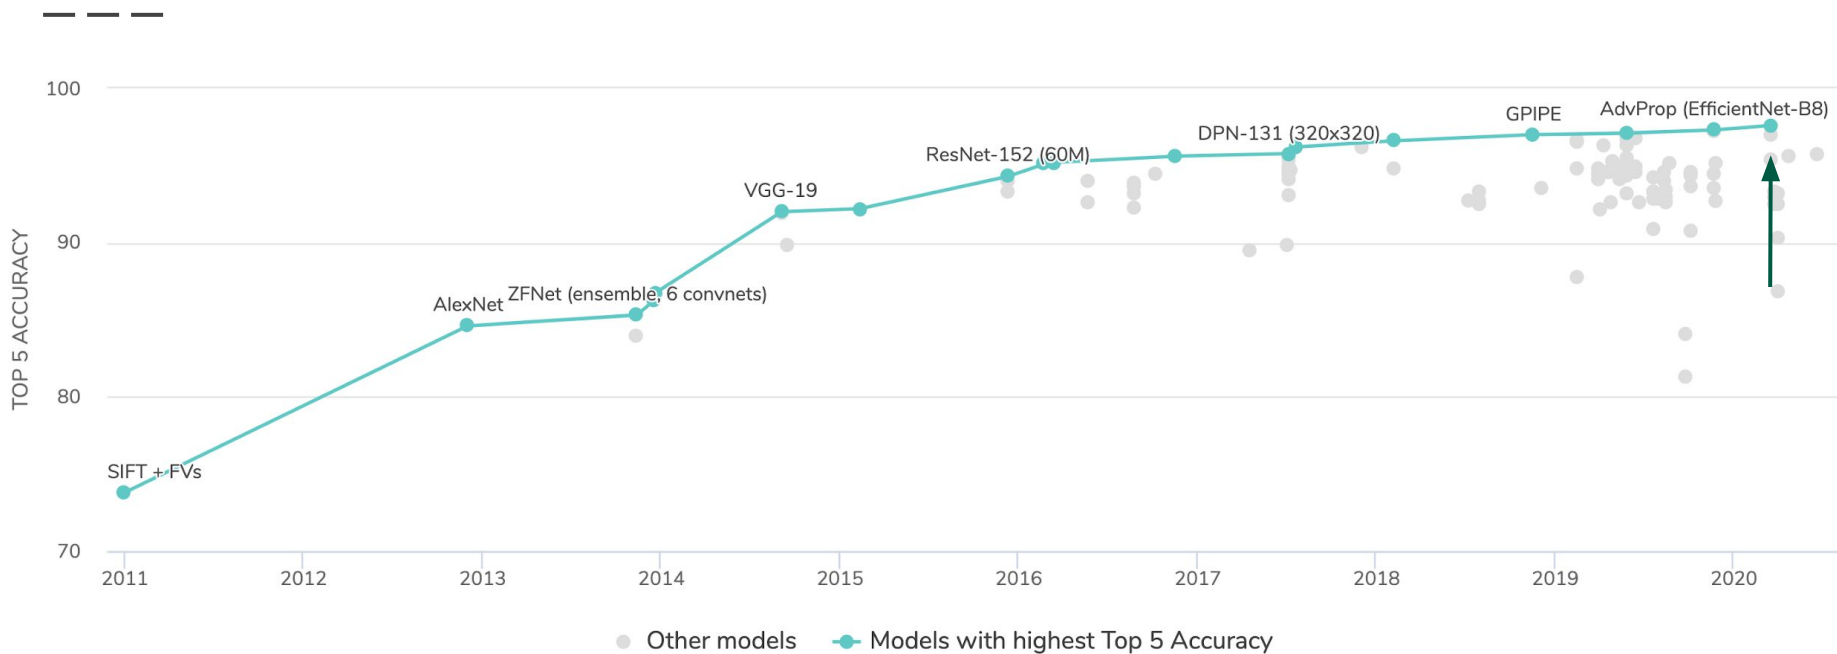
Assuming same padding  
Vanilla conv implementation  
Only counting conv multiplies

# Going deeper with the MobileNet block

— — —

- Solving for same FLOPS
- $qx(9C+2C^2) = 9C^2$
- $q = 9C / (2C+9)$
- $C = 64 \gg q \sim 4$
- Usually  $q = 6$
- ...  $q$  is an **expansion** parameter.
- “Inverted residual” or MobileNet V2
- [Sandler, Mark, et al. "Mobilenetv2: Inverted residuals and linear bottlenecks." Proceedings of the IEEE conference on computer vision and pattern recognition. 2018.](#)

# ImageNet Classification Challenge

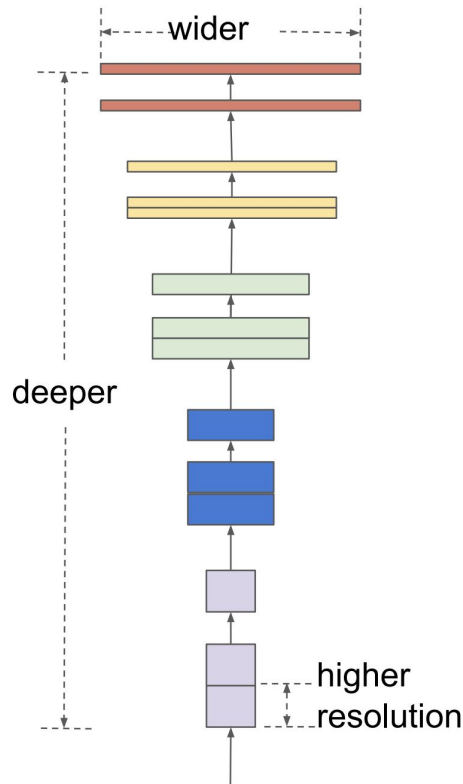


<https://paperswithcode.com/sota/image-classification-on-imagenet>

# Scaling CNNs - EfficientNet

# So many hyper-parameters...

- Number of layers = depth
- Number of channels per layer = width
- Expansion parameter of each layer =  $q$
- Resolution of each layer
- **Need to be balanced w.r.t. Compute constraint**
- **EfficientNet** = figured out a balancing using a heuristic + grid search

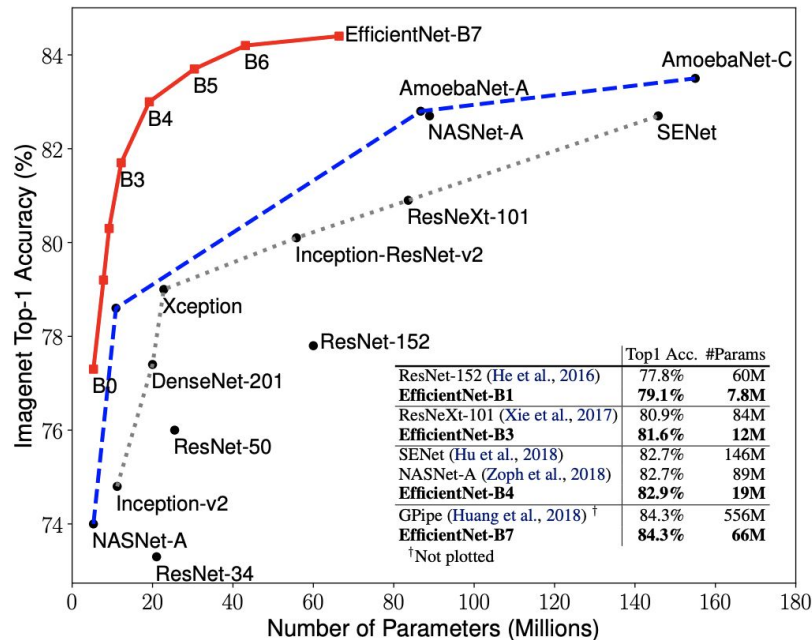


[Tan, Mingxing, and Quoc V. Le. "Efficientnet: Rethinking model scaling for convolutional neural networks." arXiv preprint arXiv:1905.11946 \(2019\).](#)

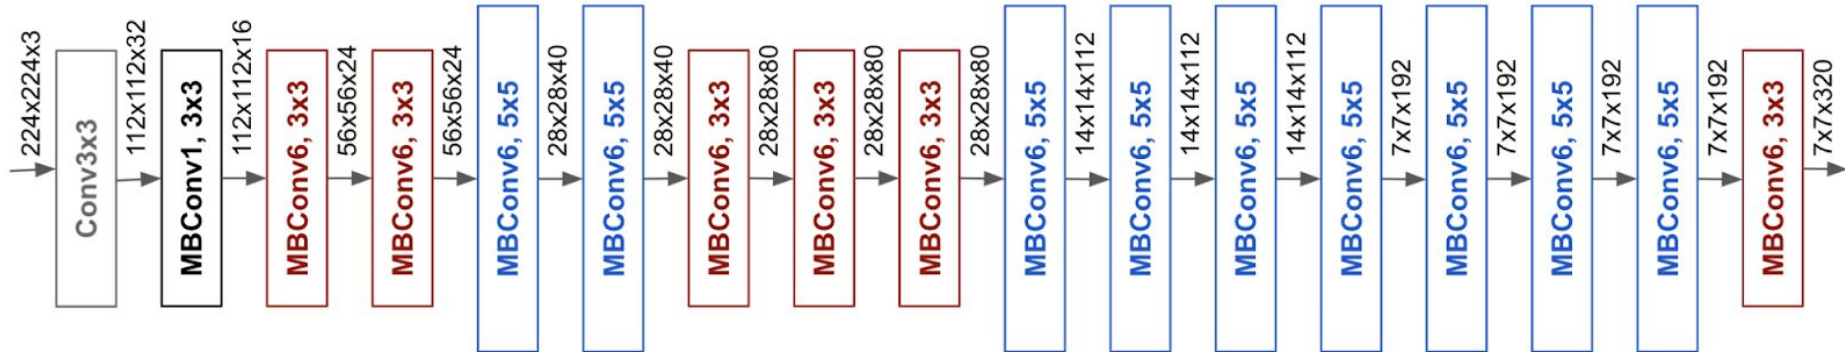


# EfficientNet

- Easy to use
- State-of-the-art (as of Oct 2020)
- Wins many kaggle competitions



# EfficientNet-B0



**End**