



Deep Learning for ECE


EECE-580G

Convolutional Neural Networks (CNNs)

From MLP to CNN


Back to MLP...

— — —



Back to MLP...

— — —



Locality


— — —

- Pixels in an image are only locally correlated (to a certain extent)
- Extends to different kinds of signals (speech, text, ...)
- Backed by biological experiments (Hubel and Wiesel's locally sensitive neurons experiments)



Locality = sparsity

— — —



Stationarity

— — —


- Patterns are shared in different locations of an image
- e.g. a vertical edge is a vertical edge no matter where in the image
- Links to translation equivariance



Credits: [Third Eye Traveler](#)

Stationarity = weight-sharing

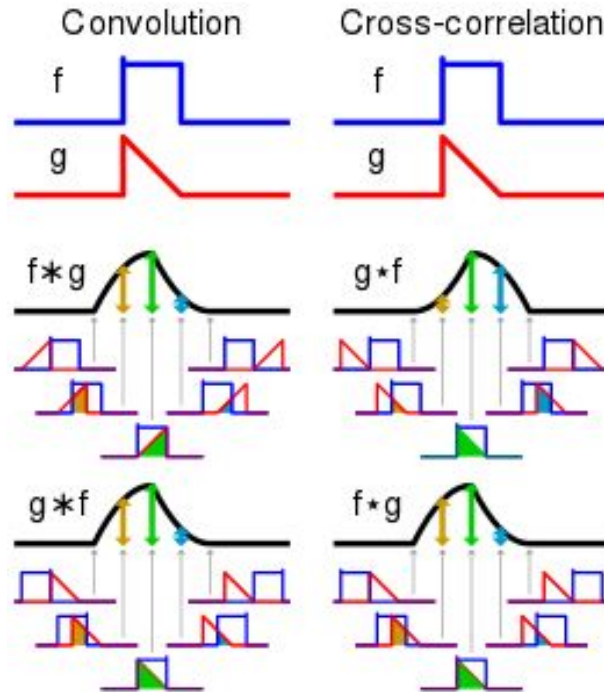
— — —



Convolution

Convolution operation

— — —



Source: [wikipedia](https://en.wikipedia.org/wiki/Convolution)


Credits

— — —

Most of the following slides will use figures created by [Vincent Dumoulin and Francesco Visin](#). Referred to as V&F.

2d convolutions

— — —



Source: V&F

2d convolutions


3_2	3_2	2_2	1	0	0
0_0	1_1	3_0	1	3	1
2_2	2_0	3_1	2	0	0
2	2	2	0	0	0
1	0	1	3	2	2
3	2	0	3	0	1

24.0	21.0	13.0	9.0
16.0	17.0	20.0	10.0
19.0	19.0	14.0	12.0
18.0	16.0	7.0	9.0

Source: V&F

2d convolutions - padding


— — —



Source: V&F

2d convolutions - strides

— — —



Source: V&F

Output shape - input shape

— — —

$$O = \lfloor I - K + 2P \rfloor / S + 1$$

O = Output size

I = Input size

K = kernel size

P = Padding

S = strides

Pooling - max

— — —


0	0	0	0	0	0	0
0	3	3	2	1	0	0
0	0	0	1	3	1	0
0	3	1	2	2	3	0
0	2	0	0	2	2	0
0	2	0	0	0	1	0
0	0	0	0	0	0	0

3.0	3.0	3.0
3.0	3.0	3.0
2.0	2.0	2.0

Source: V&F

2d convolutions - a-trous / dilation

— — —



Source: V&F

Convolutions in image processing

Box blur

— — —

$$k = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

>>> Colab

Original

— — —



Box blur - 5x5

— — —



Box blur - 11x11

— — —




Gaussian blur

— — —

$$k(i, j) = \frac{1}{2\pi\sigma^2} e^{-\frac{i^2+j^2}{2\sigma^2}}$$

>>> Colab



Original

— — —



Gaussian blur $\sigma = 2$

— — —



Gaussian blur $\sigma = 4$

— — —



Edge detectors - Sobel operator

$$k^{[1]} = \begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix}$$

$$k^{[2]} = \begin{bmatrix} +1 & +2 & +1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

>>> Colab

Edge detectors - Sobel operator

— — —



Edge detectors - Sobel operator - $k^{[1]}$

— — —



Edge detectors - Sobel operator - $k^{[2]}$

— — —



Edge detectors - Sobel operator







— — —



Stacking convolutions


How to deal with color images?

— — —




How to deal with color images?

— — —




How to deal with color images?

— — —



More filters


— — —



$A^{[i-1]}$




$K^{[i]}$




$Z^{[i]}$

Images as volumes



$A^{[i-1]}$




$K^{[i]}$




$Z^{[i]}$

Images as volumes


— — —



$A^{[i-1]}$




$K^{[i]}$




$Z^{[i]}$

Images as volumes


— — —



$A^{[i-1]}$




$K^{[i]}$




$Z^{[i]}$

Images as volumes


— — —



$A^{[i-1]}$




$K^{[i]}$




$Z^{[i]}$

Images as volumes


— — —



$A^{[i-1]}$



$K^{[i]}$



$Z^{[i]}$

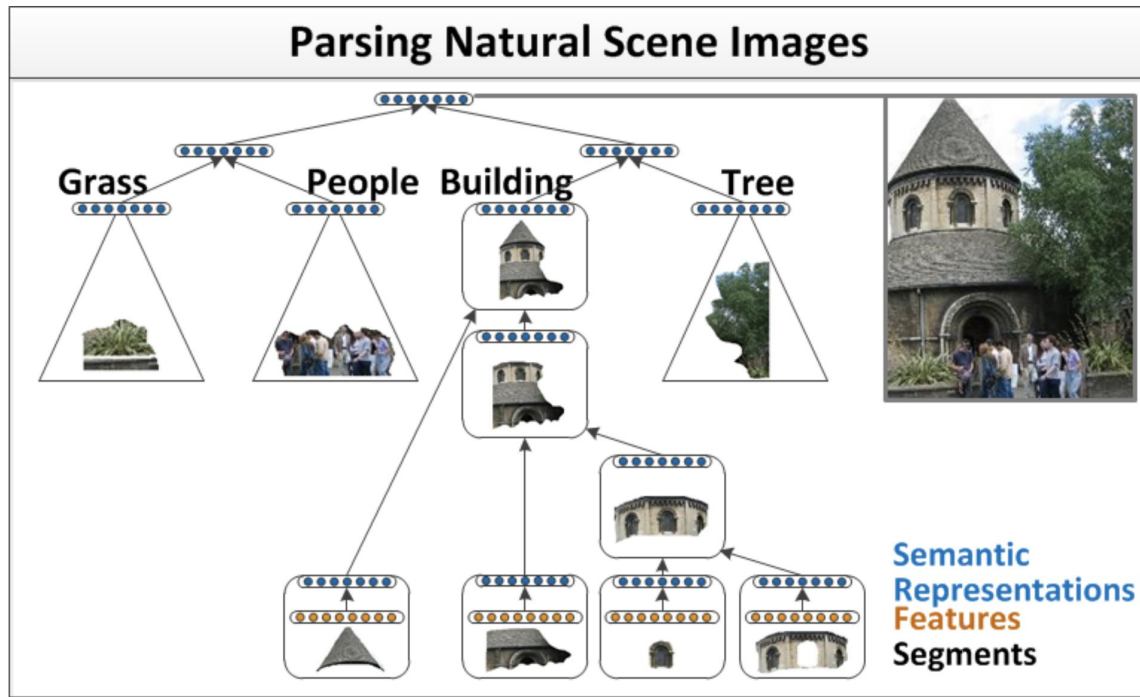
The conv kernel

— — —

$$\text{shape}(k^{[i]}) = K^2 \times \text{num channels}^{[i-1]} \times \text{num filters}$$

Compositionality

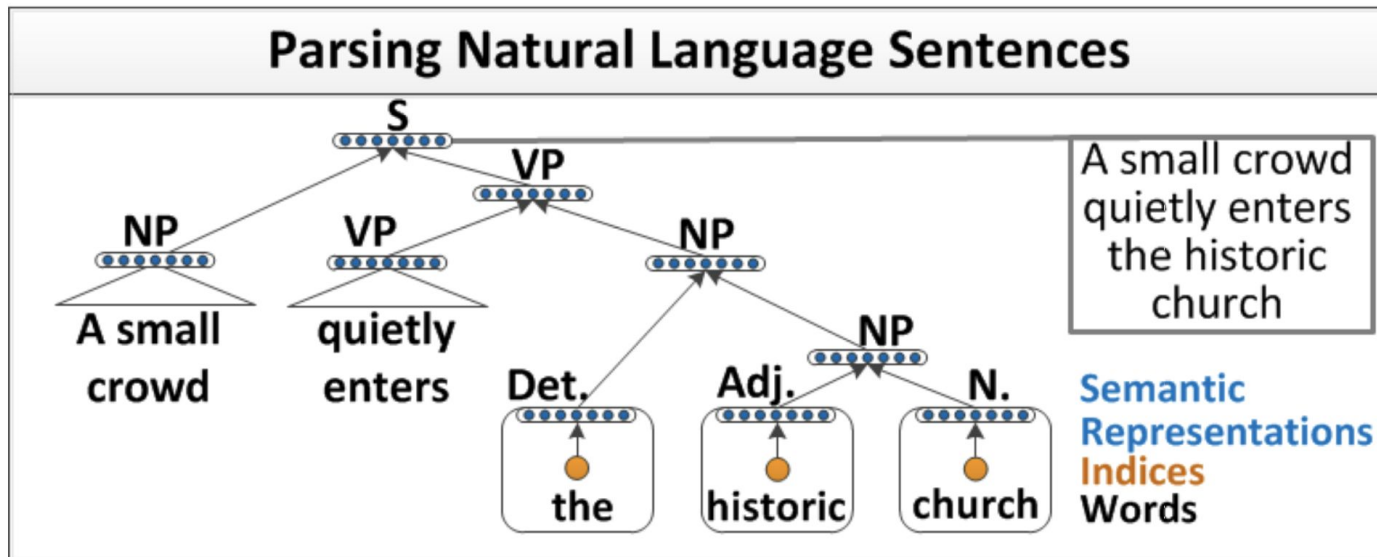
— — —



Source: [Socher, Richard, et al. "Parsing natural scenes and natural language with recursive neural networks." Proceedings of the 28th international conference on machine learning \(ICML-11\). 2011.](#)

Compositionality

— — —



Source: [Socher, Richard, et al. "Parsing natural scenes and natural language with recursive neural networks." Proceedings of the 28th international conference on machine learning \(ICML-11\). 2011.](#)

Compositionality = hierarchical representation

- Stacking multiple conv layers
- A basic conv layer = **conv** - (BN) - **activation**

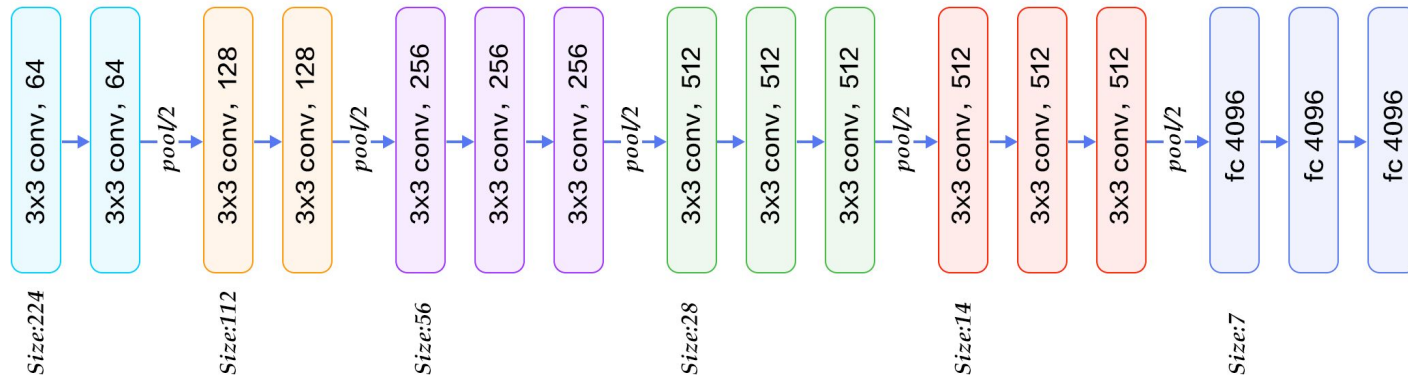



Figure Credits: Yugandhar Nanda's answer on Quora: [What is the VGG neural network?](#)
VGG paper: [Simonyan, K. and A. Zisserman. "Very Deep Convolutional Networks for Large-Scale Image Recognition." CoRR abs/1409.1556 \(2015\)](#)

Receptive field

— — —

- 3x3 conv - 3x3 conv pattern is very common
- Roughly equivalent to 5x5 conv (same receptive field)
- But 18/25 less parameters
- <https://distill.pub/2019/computing-receptive-fields/>



Pool VS Stride

— — —

- Many forensics CNNs (steganalysis, forgery detection, ...) do not use pooling in early layers
- Discarding pooling layers has also been found to be important in training good GANs
- Pooling can be criticized for losing information
- Recent CNN architectures tend to replace pooling layers with strided conv

Historical note

- The early CNNs had strides: [LeCun, Yann, et al. "Backpropagation applied to handwritten zip code recognition." Neural computation 1.4 \(1989\): 541-551.](#)
- Replaced by pooling to make the CNN more robust to pixel level distortions

[Springenberg, Jost Tobias et al. "Striving for Simplicity: The All Convolutional Net." CoRR abs/1412.6806 \(2015\):](#)

End